

```

import numpy as np
import pandas as pd
np.random.seed(42)
N = 20000 # number of users

# 1. Basic identifiers
user_id = np.arange(1, N + 1)

# 2. Demographics
age = np.clip(np.random.normal(loc=35, scale=10, size=N), 18,
70).astype(int)
countries = ["US", "UK", "IN", "CA", "DE", "BR", "AU"]
country_probs = [0.35, 0.1, 0.2, 0.1, 0.1, 0.1, 0.05]
country = np.random.choice(countries, size=N, p=country_probs)

# 3. Subscription tenure (months)
months_subscribed = np.random.randint(1, 61, size=N) # 1–60 months

# 4. Plan type
plan_types = ["basic", "standard", "premium"]
plan_probs = [0.4, 0.4, 0.2]
plan_type = np.random.choice(plan_types, size=N, p=plan_probs)

# 5. Payment method
payment_methods = ["credit_card", "paypal", "gift_card",
"bank_transfer"]
payment_probs = [0.5, 0.2, 0.1, 0.2]
payment_method = np.random.choice(payment_methods, size=N,
p=payment_probs)

# 6. Devices / profiles
# more devices for higher plans
num_devices = (
    np.where(plan_type == "basic", np.random.randint(1, 3, size=N),
    np.where(plan_type == "standard", np.random.randint(2, 4, size=N),
    np.random.randint(3, 6, size=N)))
)

# 7. Add-ons
has_4k = np.where(plan_type == "premium",
    np.random.binomial(1, 0.7, size=N),
    np.random.binomial(1, 0.05, size=N))
has_offline_downloads = np.random.binomial(1, 0.4, size=N)

# 8. Monthly price (roughly based on plan + add-ons)
base_price = np.where(plan_type == "basic", 9.99,
    np.where(plan_type == "standard", 14.99, 19.99))
monthly_price = (
    base_price
)

```

```

+ has_4k * 3.0
+ has_offline_downloads * 1.5
+ np.random.normal(0, 0.75, size=N) # small noise
)
monthly_price = np.round(monthly_price, 2)

# 9. Engagement – watch hours per last 30 days
# More expensive plans tend to watch slightly more
plan_watch_factor = np.where(plan_type == "basic", 1.0,
                             np.where(plan_type == "standard", 1.15, 1.3))
base_watch = np.random.gamma(shape=2.0, scale=8.0, size=N) # ~0–80
hours
watch_hours_30d = base_watch * plan_watch_factor
watch_hours_30d = np.clip(watch_hours_30d, 0, 120)

# 10. Engagement last 7 days – could have dropped/increased
eng_change_factor = np.random.normal(loc=0.25, scale=0.15, size=N)
# convert 30d to 7d baseline (~1/4) then apply change
watch_hours_7d = (watch_hours_30d / 4.0) * (1 + eng_change_factor)
watch_hours_7d = np.clip(watch_hours_7d, 0, None)

# 11. Days since last watch (if very low engagement, often high)
base_inactivity = np.random.randint(0, 15, size=N)
extra_inactivity = np.where(watch_hours_7d < 3, np.random.randint(0,
20,
size=N), 0)
days_since_last_watch = np.clip(base_inactivity + extra_inactivity, 0,
45)

# 12. Genre diversity & binge behaviour
genre_diversity = np.clip(
np.random.poisson(lam=3 + watch_hours_30d / 30, size=N), 1, 15
)
binge_sessions = np.clip(
    np.random.poisson(lam=watch_hours_30d / 15, size=N), 0, 20
)

# 13. Support tickets & payment failures
support_tickets = np.clip(
    np.random.poisson(lam=0.2 + (days_since_last_watch > 20) * 0.3,
size=N),
    0,
    10,
)
payment_failures = np.clip(np.random.poisson(lam=0.1 + (monthly_price
> 18) * 0.2, size=N),
    0,
    6,
)

```

```

# 14. Discount usage
discount_used = np.random.binomial(1, 0.3, size=N)

# 15. Lifetime value & engagement drop ratio
lifetime_value = monthly_price * months_subscribed
engagement_drop_ratio = np.where(
    watch_hours_30d > 0,
    watch_hours_7d / (watch_hours_30d / 4.0), # compare 7d vs
"expected" 7d
    0,
)
# clip extreme ratios
engagement_drop_ratio = np.clip(engagement_drop_ratio, 0, 2.5)

# 16. Compute churn probability using a scoring formula
# Start with a base probability
base_p = 0.15
# More likely to churn if:
p = np.full(N, base_p)
# Low watch in 30 days
p += np.where(watch_hours_30d < 10, 0.12, 0)
p += np.where(watch_hours_30d < 5, 0.10, 0)
# Big engagement drop (ratio < 0.7)
p += np.where(engagement_drop_ratio < 0.7, 0.15, 0)
# Many days since last watch
p += (days_since_last_watch / 60.0)
# Many payment failures
p += payment_failures * 0.06
# Many support tickets
p += support_tickets * 0.03
# Very new users churn more
p += np.where(months_subscribed <= 3, 0.10, 0)
# Very long-tenure users churn less
p -= np.where(months_subscribed > 24, 0.06, 0)
# Premium users churn slightly less
p -= np.where(plan_type == "premium", 0.05, 0)
# ensure probabilities are between 0.01 and 0.8
p = np.clip(p, 0.01, 0.8)

# 17. Sample actual churn outcome from Bernoulli distribution
churn = np.random.binomial(1, p, size=N)

# 18. Build DataFrame
df = pd.DataFrame({
    "user_id": user_id,
    "age": age,
    "country": country,
    "months_subscribed": months_subscribed,
    "plan_type": plan_type,
    "payment_method": payment_method,
})

```

```

    "num_devices": num_devices,
    "has_4k": has_4k,
    "has_offline_downloads": has_offline_downloads,
    "monthly_price": monthly_price,
    "watch_hours_30d": watch_hours_30d,
    "watch_hours_7d": watch_hours_7d,
    "days_since_last_watch": days_since_last_watch,
    "genre_diversity": genre_diversity,
    "binge_sessions": binge_sessions,
    "support_tickets": support_tickets,
    "payment_failures": payment_failures,
    "discount_used": discount_used,
    "lifetime_value": lifetime_value,
    "engagement_drop_ratio": engagement_drop_ratio,
    "churn": churn
})

# 19. Save to CSV
df.to_csv("streaming_churn_20k.csv", index=False)
print("Saved dataset with shape:", df.shape)
print(df.head())

```

Saved dataset with shape: (20000, 21)

	user_id	age	country	months_subscribed	plan_type	payment_method	...
0	1	39	US		12	basic	credit_card
1	2	33	AU		58	standard	bank_transfer
2	3	41	DE		49	basic	credit_card
3	4	50	BR		34	standard	paypal
4	5	32	CA		49	standard	paypal

	num_devices	has_4k	has_offline_downloads	monthly_price	...
0	1	0	1	12.06	...
1	2	1	0	18.44	...
2	1	0	1	10.16	...
3	2	0	1	16.52	...
4	2	0	1	16.79	...

	watch_hours_7d	days_since_last_watch	genre_diversity
binge_sessions	6.818861	1	7
0	5.046324	0	5
2	5.446575	4	3
1	4.044258	7	3
2	6.162215	11	1
0			

```

support_tickets payment_failures discount_used lifetime_value \
0           1                 0             1       144.72
1           0                 1             1      1069.52
2           0                 0             0       497.84
3           1                 0             0       561.68
4           0                 0             0       822.71

engagement_drop_ratio  churn
0        1.331173     0
1        1.211439     0
2        1.072884     0
3        1.297427     0
4        1.421602     0

[5 rows x 21 columns]

#20 EDA BASIC CHECKING
import pandas as pd

df = pd.read_csv("streaming_churn_20k.csv")
df.info()
df.describe(include="all")
df["churn"].value_counts(normalize=True)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 21 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   user_id           20000 non-null   int64  
 1   age               20000 non-null   int64  
 2   country           20000 non-null   object  
 3   months_subscribed 20000 non-null   int64  
 4   plan_type          20000 non-null   object  
 5   payment_method     20000 non-null   object  
 6   num_devices        20000 non-null   int64  
 7   has_4k             20000 non-null   int64  
 8   has_offline_downloads 20000 non-null   int64  
 9   monthly_price      20000 non-null   float64
 10  watch_hours_30d    20000 non-null   float64
 11  watch_hours_7d     20000 non-null   float64
 12  days_since_last_watch 20000 non-null   int64  
 13  genre_diversity   20000 non-null   int64  
 14  binge_sessions     20000 non-null   int64  
 15  support_tickets    20000 non-null   int64  
 16  payment_failures   20000 non-null   int64  
 17  discount_used      20000 non-null   int64  
 18  lifetime_value     20000 non-null   float64
 19  engagement_drop_ratio 20000 non-null   float64
 20  churn              20000 non-null   int64

```

```
dtypes: float64(5), int64(13), object(3)
memory usage: 3.2+ MB

churn
0      0.6663
1      0.3337
Name: proportion, dtype: float64


#21 UNIVARIATE EDA
import seaborn as sns
import matplotlib.pyplot as plt


num_cols =
["age","months_subscribed","monthly_price","watch_hours_30d",
 "watch_hours_7d","days_since_last_watch","lifetime_value",


"engagement_drop_ratio","payment_failures","support_tickets"]
cat_cols = ["country","plan_type","payment_method","churn"]


for col in num_cols:
    sns.histplot(df[col], kde=True)
    plt.title(col)
    plt.show()

for col in cat_cols:
    sns.countplot(data=df, x=col)
    plt.title(col)
    plt.show()
```














watch_hours_7d


























```
#22 SINGLE VS CHURN
# Churn rate by plan
print(df.groupby("plan_type")["churn"].mean())

# Churn rate by country
print(df.groupby("country")["churn"].mean())

# Churn vs engagement buckets
df["watch_bucket"] = pd.cut(df["watch_hours_30d"],
bins=[0,10,30,60,120],
labels=["0-10","10-30","30-60","60-120"])
print(df.groupby("watch_bucket")["churn"].mean())


plan_type
basic      0.352223
premium    0.292907
standard   0.336224
Name: churn, dtype: float64
country
AU      0.350663
BR      0.349391
CA      0.331994
DE      0.342119
```


```
IN      0.325552
UK      0.322368
US      0.332865
Name: churn, dtype: float64
watch_bucket
0-10     0.539014
10-30    0.245915
30-60    0.235131
60-120   0.204762
Name: churn, dtype: float64

C:\Users\rithw\AppData\Local\Temp\ipykernel_1064\1764451689.py:11:
FutureWarning: The default of observed=False is deprecated and will be
changed to True in a future version of pandas. Pass observed=False to
retain current behavior or observed=True to adopt the future default
and silence this warning.
print(df.groupby("watch_bucket")["churn"].mean())

sns.barplot(data=df, x="plan_type", y="churn")
plt.show()

sns.boxplot(data=df, x="churn", y="watch_hours_30d")
plt.show()
```






#23 CORRELATIONS

```
corr = df.corr(numeric_only=True)
sns.heatmap(corr, cmap="coolwarm", center=0)
plt.show()
```

```
corr["churn"].sort_values(ascending=False)
```



churn	1.000000
days_since_last_watch	0.329457
support_tickets	0.080185
payment_failures	0.026488
user_id	0.017029
discount_used	0.003831
age	0.003822
has_offline_downloads	0.000081
engagement_drop_ratio	-0.018274
has_4k	-0.039108
num_devices	-0.043888
genre_diversity	-0.044630
monthly_price	-0.045250
months_subscribed	-0.072744
lifetime_value	-0.086147
binge_sessions	-0.136070
watch_hours_30d	-0.195112

```
watch_hours_7d      -0.195669  
Name: churn, dtype: float64
```